DistilBERT on SST-2: Reproduction Brief

Brief restating results from Sanh, Debut, Chaumond and Wolf (2019), "DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter", arXiv:1910.01108. It is not the original paper; see the original for full details.

1. Method

DistilBERT is a 6-layer Transformer student distilled from BERT-base during pre-training with a triple loss (distillation, masked language modeling and cosine embedding). For SST-2 the model is fine-tuned as a binary sentence classifier and evaluated on the 872-sentence development set of the GLUE benchmark.

2. Results

On the SST-2 development set DistilBERT reaches an accuracy of 91.3%, compared with 92.7% for BERT-base (Table 1 of the original paper; GLUE dev-set results are medians of 5 runs).

3. Released checkpoint

The fine-tuned DistilBERT SST-2 checkpoint is released on the Hugging Face Hub at https://huggingface.co/distilbert/distilbert-base-uncased-finetuned-sst-2-english